

Saddle-Dominated Capacity of the LSE Dense Associative Memory

Abstract

[Placeholder.] The log-sum-exp Dense Associative Memory [3] stores M unit vectors ξ^μ on the N -sphere with kernel sharpness β_{net} . Ramsauer et al. proved $\alpha_c = 1/2$ at $T = 0$ by modelling the inter-pattern overlaps as Gaussian. Replacing the Gaussian by the exact spherical density opens an interior saddle in the partition sum at $\varphi^*(\beta_{\text{net}})$, with value $g_{\text{max}}(\beta_{\text{net}})$. The boundary becomes $\alpha_c(\beta_{\text{net}}, T) = \beta_{\text{net}}[1 - f_{\text{ret}}(T)] - g_{\text{max}}(\beta_{\text{net}})$. At $\beta_{\text{net}} = 1$ — the value we use in the Monte Carlo — the saddle is the golden ratio $\varphi^* = (\sqrt{5} - 1)/2$ and $g_{\text{max}} \approx 0.3774$, giving $\alpha_c(0) \approx 0.6226$, not 0.5. Ramsauer’s $1/2$ is β_{net} -independent (an extreme-value cancellation), the exact-density boundary is not. We confirm the corrected value with three Monte Carlo variants (direct, truncated, hybrid) at N up to 100.

1 Introduction

A Dense Associative Memory [1] stores M patterns $\xi^\mu \in \mathbb{R}^N$, $\mu = 1, \dots, M$, as minima of an energy function $H(\mathbf{x})$. Retrieval of ξ^1 from a noisy cue is the Metropolis dynamics of $\mathbf{x}$ under H at external temperature T . Both the patterns and the chain live on the N -sphere of radius $\sqrt{N}$. The log-sum-exp (LSE) variant of Ramsauer et al. uses

$$H(\mathbf{x}) = -\beta_{\text{net}}^{-1} \ln \sum_{\mu=1}^M \exp(\beta_{\text{net}} \mathbf{x} \cdot \xi^\mu), \quad \beta_{\text{net}} = 1 \text{ throughout}, \quad (1)$$

and stores $M = e^{\alpha N}$ patterns at load $\alpha = \ln M/N$. The question is: what is the largest α at which the chain initialised at ξ^1 stays close to ξ^1 ?

Ramsauer et al. answered $\alpha_c = 1/2$ at $T = 0$. The proof modelled the inter-pattern overlap $\varphi_{1\mu} = \xi^1 \cdot \xi^\mu / N$ as a Gaussian with variance $1/N$. The Gaussian is the leading term of the exact spherical density

$$f(\varphi) \propto (1 - \varphi^2)^{(N-3)/2}, \quad \varphi \in [-1, 1], \quad (2)$$

near $\varphi = 0$, but the two densities differ near $\varphi = 1$. That tail matters: it sets the typical maximum of M samples and the saddle of the partition sum.

This paper does three things. First, we re-derive the boundary using the exact density (2). An interior saddle of the integrand appears at a point $\varphi^*(\beta_{\text{net}})$ that depends on the kernel sharpness; at $\beta_{\text{net}} = 1$ — the value we use throughout — it is the golden ratio $\varphi^* = (\sqrt{5} - 1)/2$, with saddle value $g_{\text{max}} \approx 0.3774$, and the boundary shifts to $\alpha_c(0) = 1 - g_{\text{max}} \approx 0.6226$. Second, we describe three Monte Carlo schemes (direct, truncated, hybrid) and explain why a naive “smart” variant that refreshes its pattern cache fails near α_c . Third, we confirm the corrected boundary numerically at $N = 100$ by an escape-time measurement and by a ϕ residual heatmap.

2 Theory

Free energy split. For a chain in the basin of ξ^1 , the LSE sum is dominated by the $\mu = 1$ term, $H(\mathbf{x})/N \approx -\varphi_1$ with $\varphi_1 = \xi^1 \cdot \mathbf{x}/N$. Define the energy per spin relative to the retrieval

baseline, $u(\varphi_1) = 1 - \varphi_1$, and the $(N-1)$ -sphere entropy per spin $s(\varphi_1) = \frac{1}{2} \ln(1 - \varphi_1^2)$. The retrieval free energy is the minimum of $f(\varphi_1, T) = u - Ts$:

$$f_{\text{ret}}(T) = (1 - \phi_{\text{eq}}) - \frac{T}{2} \ln(1 - \phi_{\text{eq}}^2), \quad \phi_{\text{eq}}(T) = \frac{1}{2}(-T + \sqrt{T^2 + 4}). \quad (3)$$

For a chain uncorrelated with ξ^1 ($\varphi_1 \approx 0$) the sum in (1) is dominated by the $M-1$ spurious terms $S = \sum_{\mu \neq 1} \exp(\beta_{\text{net}} N \varphi_{1\mu})$. Replacing the sum by an integral with density (2) and using $M = e^{\alpha N}$:

$$S \approx e^{\alpha N} \int_{-1}^1 \exp\left[N\left(\frac{1}{2} \ln(1 - \varphi^2) + \beta_{\text{net}} \varphi\right)\right] d\varphi \equiv e^{\alpha N} \int e^{Ng(\varphi; \beta_{\text{net}})} d\varphi. \quad (4)$$

$g(\varphi; \beta_{\text{net}}) = \frac{1}{2} \ln(1 - \varphi^2) + \beta_{\text{net}} \varphi$ has an interior saddle at $g'(\varphi^*) = 0$, i.e. $\beta_{\text{net}}(1 - \varphi^{*2}) = \varphi^*$, with positive root

$$\varphi^*(\beta_{\text{net}}) = \frac{-1 + \sqrt{1 + 4\beta_{\text{net}}^2}}{2\beta_{\text{net}}}. \quad (5)$$

The saddle value is

$$g_{\text{max}}(\beta_{\text{net}}) = g(\varphi^*; \beta_{\text{net}}) = \frac{1}{2} \ln(1 - \varphi^{*2}) + \beta_{\text{net}} \varphi^*. \quad (6)$$

Hence $N^{-1} \ln S \rightarrow \alpha + g_{\text{max}}(\beta_{\text{net}})$, the noise-floor energy per spin is $u_{\text{noise}} = 1 - [\alpha + g_{\text{max}}(\beta_{\text{net}})]/\beta_{\text{net}}$, and the boundary at the crossover $u_{\text{noise}} = f_{\text{ret}}(T)$ is

$$\boxed{\alpha_c(\beta_{\text{net}}, T) = \beta_{\text{net}} [1 - f_{\text{ret}}(T)] - g_{\text{max}}(\beta_{\text{net}}).} \quad (7)$$

Gaussian and support-edge variants. Two well-known prior formulas come from substituting different objects into (4) and using its *support edge* instead of the interior saddle.

Gaussian density. Replace $f(\varphi)$ by the moment-matched Gaussian $\sqrt{N/2\pi} \exp(-N\varphi^2/2)$. The integrand exponent becomes $g_G(\varphi; \beta_{\text{net}}) = -\frac{1}{2}\varphi^2 + \beta_{\text{net}}\varphi$, monotone on $[-1, 1]$ for $\beta_{\text{net}} \leq 1$. The integral is dominated by the typical maximum of M Gaussian samples $\varphi_{\text{max}}^{(G)}(\alpha) = \sqrt{2\alpha}$; at this edge $g_G(\varphi_{\text{max}}^{(G)}) = \beta_{\text{net}}\sqrt{2\alpha} - \alpha$, so $N^{-1} \ln S \rightarrow \beta_{\text{net}}\sqrt{2\alpha}$ and $u_{\text{noise}}^{(G)} = 1 - \sqrt{2\alpha}$. Solving $u_{\text{noise}}^{(G)} = f_{\text{ret}}(T)$:

$$\alpha_c^G(T) = \frac{1}{2}[1 - f_{\text{ret}}(T)]^2, \quad \alpha_c^G(0) = \frac{1}{2}. \quad (8)$$

β_{net} cancels — both retrieval and spurious-extreme energies are $N \cdot \beta_{\text{net}}$ times an overlap, and the β_{net}^{-1} prefactor in H removes it. This is Ramsauer's threshold.

Exact density, support edge. Keep $f(\varphi)$ but use the support edge of the exact density $\varphi_{\text{max}}(\alpha) = \sqrt{1 - e^{-2\alpha}}$ (from $M \cdot \Pr_f(\varphi > \varphi_{\text{max}}) = 1$). At this edge $\frac{1}{2} \ln(1 - \varphi_{\text{max}}^2) = -\alpha$, so $g(\varphi_{\text{max}}) = -\alpha + \beta_{\text{net}} \varphi_{\text{max}}$, $N^{-1} \ln S \rightarrow \beta_{\text{net}} \varphi_{\text{max}}(\alpha)$ and $u_{\text{noise}}^{(\text{ex})} = 1 - \varphi_{\text{max}}(\alpha)$. Solving $1 - \varphi_{\text{max}}(\alpha_c^{\text{bd}}) = f_{\text{ret}}(T)$:

$$\alpha_c^{\text{bd}}(T) = -\frac{1}{2} \ln[1 - (1 - f_{\text{ret}}(T))^2], \quad \alpha_c^{\text{bd}}(0) = +\infty. \quad (9)$$

This is the formula that works for the LSR kernel where g is monotone on the support [2]; for LSE it is correct only when the interior saddle φ^* lies outside the support $[0, \varphi_{\text{max}}(\alpha)]$, i.e. for $\alpha < -\frac{1}{2} \ln(1 - \varphi^{*2}) \approx 0.241$ at $\beta_{\text{net}} = 1$. Above this threshold the interior saddle (7) takes over. Both (8) and (9) are β_{net} -independent for the same reason: the support edge replaces the β_{net} -dependent saddle.

Cusp picture. At fixed φ_1 the spurious sum evaluates to $e^{N(\alpha + g_{\max}(\beta_{\text{net}}))}$ to leading order, independent of φ_1 . Inside the LSE log-sum, $\beta_{\text{net}}^{-1} \ln[e^{\beta_{\text{net}} N \varphi_1} + S] \rightarrow \max(\varphi_1, [\alpha + g_{\max}(\beta_{\text{net}})]/\beta_{\text{net}})$ at large N , so the leading-order free energy per spin reads

$$\frac{F(\varphi_1)}{N} = -\max\left(\varphi_1, \frac{\alpha + g_{\max}(\beta_{\text{net}})}{\beta_{\text{net}}}\right) - \frac{T}{2} \ln(1 - \varphi_1^2), \quad (10)$$

with a kink (“cusp”) at $\varphi_c(\alpha, \beta_{\text{net}}) = [\alpha + g_{\max}(\beta_{\text{net}})]/\beta_{\text{net}}$. As T rises the basin minimum $\phi_{\text{eq}}(T)$ slides toward the cusp; at $\phi_{\text{eq}}(T) = \varphi_c$ the local minimum disappears, defining the single-pattern spinodal

$$\alpha_c^{\text{sp}}(\beta_{\text{net}}, T) = \beta_{\text{net}} \phi_{\text{eq}}(T) - g_{\max}(\beta_{\text{net}}), \quad T_{\text{sp}}(\alpha, \beta_{\text{net}}) = \frac{\beta_{\text{net}}^2 - (\alpha + g_{\max}(\beta_{\text{net}}))^2}{\beta_{\text{net}} (\alpha + g_{\max}(\beta_{\text{net}}))}. \quad (11)$$

$\alpha_c^{\text{sd}}(\beta_{\text{net}}, T)$ and $\alpha_c^{\text{sp}}(\beta_{\text{net}}, T)$ meet at the endpoint $(\beta_{\text{net}} - g_{\max}(\beta_{\text{net}}), 0)$.

Examples and the $\beta_{\text{net}} \rightarrow \infty$ limit. Four cases of (5), (6), and (7) at $T = 0$:

β_{net}	$\varphi^*(\beta_{\text{net}})$	$g_{\max}(\beta_{\text{net}})$	$\alpha_c(\beta_{\text{net}}, 0)$	note
0.5	$\sqrt{2} - 1 \approx 0.414$	≈ 0.113	≈ 0.387	—
1	$(\sqrt{5} - 1)/2 \approx 0.618$ (golden ratio)	≈ 0.3774	≈ 0.6226	MC value
2	$(\sqrt{17} - 1)/4 \approx 0.781$	≈ 1.091	≈ 0.909	—
$\rightarrow \infty$	$\rightarrow 1 - 1/(2\beta_{\text{net}})$	$\beta_{\text{net}} - \frac{1}{2}(\ln \beta_{\text{net}} + 1)$	$\frac{1}{2}(\ln \beta_{\text{net}} + 1)$	log growth

Capacity grows logarithmically with β_{net} . The Ramsauer prediction $\alpha_c^{\text{G}} = 1/2$ is β_{net} -independent: its derivation replaces the integrand in (4) by its support edge $\sqrt{2\alpha}$, and the factor β_{net} then cancels between retrieval and spurious terms. The interior saddle restores the β_{net} dependence; the β_{net} -independence of the Gaussian-edge answer is itself the signature of the support-edge replacement. The MC numerics in Sections 3 and 4 fix $\beta_{\text{net}} = 1$, where $\alpha_c(1, 0) \approx 0.6226$ is the value we verify. Figure 1 illustrates (10) at $\alpha = 0.50$, $\beta_{\text{net}} = 1$.

3 Monte Carlo schemes

All chains run on the $\sqrt{N}$ -sphere at $\beta_{\text{net}} = 1$ with Metropolis updates. The variants differ in how they handle the $M = e^{\alpha N}$ patterns.

Direct. Store all M patterns explicitly and evaluate (1) over all of them at every step. Memory and compute cost $\sim MN$. Feasible up to $N \approx 25$ at $\alpha \approx 0.7$ ($M \approx 4 \times 10^7$); impossible at $N = 100$, $\alpha = 0.6$ ($M \approx 10^{26}$).

Truncated. Generate only the patterns with overlap $\varphi_{1\mu} > \varphi_{\text{keep}}$ to the target ξ^1 . The cutoff is chosen so the expected count above it equals a budget $K \sim 10^4$; the actual count is Poisson with mean $M \cdot \Pr(\varphi > \varphi_{\text{keep}})$, and the K overlap values are drawn from the Beta tail. The patterns below φ_{keep} are not generated; their contribution enters the LSE log-sum as the analytic constant $\ln C_{\text{bulk}} = \alpha N + \ln \int_{-1}^{\varphi_{\text{keep}}} f(\varphi) e^{N\varphi} d\varphi$. The retained set is fixed once per disorder. Truncated MC at $N = 25$ matches direct MC on the basin boundary across the full α range we measure, so the $\varphi < \varphi_{\text{keep}}$ patterns are genuinely unnecessary for the basin physics, not silently dropped. We use truncated for the wide heatmap.

Hybrid. Same kept-tail + analytic-bulk decomposition as truncated, but the chain is reduced to its scalar overlap φ_1 and the bulk constant is set by the saddle of (4). Concretely: patterns with overlap $\varphi_{1\mu} > \varphi_c = [\alpha + g_{\max}(\beta_{\text{net}})]/\beta_{\text{net}}$ are sampled from the spherical Beta tail and kept fixed per disorder (count $K_{\text{cusp}} \sim M \cdot \Pr(\varphi > \varphi_c)$; in our $N = 100$, $\beta_{\text{net}} = 1$ runs, between




Figure 1: Leading-order free energy per spin $F(\varphi_1)/N$ from (10) at $\alpha = 0.50$, $\beta_{\text{net}} = 1$, for several T . The cusp at $\varphi_c = \alpha + g_{\text{max}} \approx 0.877$ is the boundary between the retrieval branch (slope -1) and the noise-floor branch (slope 0). Dots: basin minima $\phi_{\text{eq}}(T)$. At $T = T_{\text{sp}}(0.50, 1) \approx 0.262$ the minimum reaches the cusp; above this T the basin no longer exists.

~ 4500 at $\alpha = 0.50$ and ~ 8500 at $\alpha = 0.62$). The contribution of the bulk ($\varphi_{1\mu} < \varphi_c$) is replaced by the φ_1 -independent constant $M e^{Ng_{\text{max}}(\beta_{\text{net}})}$ from (6). Metropolis runs in 1D on φ_1 with the $(N-1)$ -sphere entropy $\frac{1}{2} \ln(1 - \varphi_1^2)$ as the prior. The physics is the same as truncated; the gain is one scalar update per MC step instead of an N -dimensional spin update, which makes the scheme cheap enough to push to $N = 100$ near α_c where truncated MC at the full spin becomes slow.

Why naive “smart” fails. A faster variant — “smart” — refreshes its K kept patterns every N_{REFRESH} steps from a cone around the current chain. This breaks the quenched-disorder assumption. As the chain wanders, fresh patterns are minted at every refresh; the total set of distinct patterns seen along one trajectory exceeds M , and the chain effectively faces a larger pattern set than the model defines. Near the retrieval basin the extra competitors destroy the basin. Both truncated and hybrid avoid the problem: the retained set is fixed once per disorder, no refresh.

4 Results

We measure two observables at $N = 100$: the mean first-passage escape time τ from the retrieval basin, and the steady-state $\langle \varphi_1 \rangle - \phi_{\text{eq}}(T)$ residual.

Escape time. Figure 2 reports $\log_{10} \langle \tau \rangle$ vs T at $N = 100$ for five α values from the N-dim hybrid MC, averaged over 32 disorders, with a budget cap at 10^6 MC steps. Each curve is a basin escape profile: at low T the basin is stable for the full budget (upward triangles at the ceiling); above a critical $T_{\text{esc}}(\alpha)$ the escape time drops by several orders of magnitude over a T window of width ~ 0.01 . T_{esc} moves to lower T as α rises; extrapolating to $\tau \rightarrow \infty$ at $T = 0$ gives a critical load near $\alpha = 1 - g_{\text{max}} \approx 0.6226$.

Residual heatmap. Figure 3 shows $\langle \varphi_1 \rangle - \phi_{\text{eq}}(T)$ over the (α, T) plane. Background: truncated MC at varying N over $\alpha \in [0.20, 0.70]$. Grey rectangle: N-dim hybrid MC at $N = 100$ on




Figure 2: Mean basin escape time $\langle \tau \rangle$ vs T at $N = 100$ from N-dim hybrid MC, 32 disorders, $\beta_{\text{net}} = 1$. Five α values: 0.50, 0.55, 0.58, 0.60, 0.62. Filled circles: mean over disorders that escaped within 10^6 MC steps. Upward triangles at $\log_{10} \tau = 6$: cells where no disorder escaped (lower bound only). The escape curve slides toward $T = 0$ as $\alpha \rightarrow 1 - g_{\text{max}} \approx 0.6226$.

a fine T grid in $\alpha \in [0.50, 0.62]$, $T \in [0, 0.10]$. The basin (white) is bounded by a sharp transition (dark red) that follows the saddle line $\alpha_c^{\text{sd}}(T)$ and pinches at $\alpha \approx 0.6226$, $T = 0$.

5 Discussion

The Ramsauer threshold $\alpha_c = 1/2$ is a Gaussian artefact. The Gaussian density misses the interior saddle of the partition-sum integrand (4) that the exact spherical density places at $\varphi^*(\beta_{\text{net}})$. Including it gives $\alpha_c(\beta_{\text{net}}, 0) = \beta_{\text{net}} - g_{\text{max}}(\beta_{\text{net}})$; at the value $\beta_{\text{net}} = 1$ used in our MC, $\alpha_c \approx 0.6226$, a factor 1.25 larger than Ramsauer’s $1/2$. The shift is per spin, not asymptotic: it appears at leading order in N and is visible already at $N = 100$. A second signature of the artefact is that Ramsauer’s $1/2$ is β_{net} -independent, while the exact-density (7) boundary is not.

Open items. (i) Hot init. All present runs start the chain at ξ^1 (cold). The basin boundary they trace is the spinodal, not the thermodynamic crossover; hot init at the same grid would delimit the metastability band. (ii) Cusp rounding at finite N . The leading-order $\max(\cdot, \cdot)$ in (10) is rounded on a scale $1/N$ in φ_1 . Near the endpoint $(0.6226, 0)$ the basin offset $\phi_{\text{eq}}(T) - (\alpha + g_{\text{max}})$ falls inside the rounding window, and the leading-order spinodal slope needs an NLO correction. (iii) Heatmap consolidation. Figure 3 stitches three data sources. A clean version would use a single MC variant — truncated with a hybrid insert near the cusp — across the entire (α, T) plane.

References

- [1] Dmitry Krotov and John J. Hopfield. Dense associative memory for pattern recognition. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 29, pages 1172–1180, 2016.
- [2] Tatiana Petrova, Evgeny Polyachenko, and Radu State. Retrieval metastability via Kramers




Figure 3: $\langle \varphi_1 \rangle - \phi_{\text{eq}}(T)$ from Monte Carlo over the (α, T) plane. Curves: $\alpha_c^G(T)$ (blue solid, Eq. 8, Gaussian density + support edge); $\alpha_c^{\text{bd}}(T)$ (orange dashed, Eq. 9, exact density + support edge); $\alpha_c^{\text{sd}}(T)$ (red solid, Eq. 7, exact density + interior saddle, this work); $\alpha_c^{\text{sp}}(T)$ (purple dotted, Eq. 11, single-pattern spinodal). Background: truncated MC at varying N over the full α range, linearly interpolated to the fine T grid. Grey rectangle: N-dim hybrid MC at $N = 100$ on the fine T grid. The basin pinches at $\alpha \approx 0.6226$ — Ramsauer’s $\alpha_c^G(0) = 1/2$ undershoots by 0.12 in α , a factor 1.25 in capacity.

escape and centroid hopping. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2026. placeholder key; replace at submission.

- [3] Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Lukas Gruber, Markus Holzleitner, Thomas Adler, David Kreil, Michael K. Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. Hopfield networks is all you need. In *International Conference on Learning Representations*, 2021.