

Thermodynamic approach and Saddle-Dominated Capacity of the LSE Dense Associative Memory

Anonymous submission

Abstract

The Dense Associative Memory with log-sum-exp (LSE) kernel on the N -sphere stores $M = \exp(\alpha N)$ patterns at load α . A Gaussian approximation of the inter-pattern alignments gives a zero-temperature capacity $\alpha_c = 1/2$, independent of the kernel sharpness β , the Gaussian-ensemble threshold of Lucibello and Mézard (2024). Using the exact spherical alignment density, we build a finite-temperature theory of retrieval. The sum over non-target patterns is then set by an interior saddle rather than by the right tail of the Gaussian, which restores the β -dependence: at $T = 0$ the capacity grows from $\sim \beta$ at small β to $\frac{1}{2} \ln \beta$ at large β , is ≈ 0.6226 at $\beta = 1$, and reproduces the exact spherical threshold of Lucibello and Mézard. At $T > 0$ the retrieval basin is bounded by a single-pattern spinodal, beyond which it no longer exists; below it the basin is metastable, with a Kramers escape time exponential in N . We confirm the resulting (α, T) stability diagram by Monte Carlo: a direct sampler stores all M patterns but reaches only $N \sim 20$, while a truncated sampler that keeps only the boundary-setting patterns reaches $N \sim 30$ and agrees with both the direct sampler and the predicted spinodal.

Introduction

Modern dense associative memory (DAM) models begin with Krotov and Hopfield (2016), who generalized the classical Hopfield network (Hopfield 1982) from pairwise to polynomial interactions, raising the storage capacity from linear to polynomial in the number of neurons N while preserving the ability to retrieve previously stored information from noisy or corrupted input data. Taking the polynomial degree to infinity yields an exponential capacity (Demircigil et al. 2017): for binary neurons (state space $\{-1, +1\}^N$), the limiting capacity is $M_c = \exp(\alpha_c N)$ with load $\alpha_c = (\ln 2)/2 \approx 0.35$.

The same exponential scaling carries over to continuous states. Ramsauer et al. (2021) introduced a modern Hopfield network whose state is a continuous vector in $\mathbb{R}^N$ and whose energy is a log-sum-exp (LSE) of the state-pattern overlaps. We take both the stored patterns ξ^μ and the retrieval state $\mathbf{x}$ on the sphere $|\mathbf{x}|^2 = N$, and store $M = \exp(\alpha N)$ patterns at load α . The energy then reads

$$H(\mathbf{x}) = -\beta^{-1} \ln \sum_{\mu=1}^M \exp[-\beta N(1 - \phi^\mu)], \quad (1)$$

where β is a kernel sharpness controlling how strongly the energy distinguishes patterns, and ϕ^μ is the *alignment* of the current state $\mathbf{x}$ with the μ -th pattern ξ^μ ,

$$\phi^\mu(\mathbf{x}) = N^{-1} \mathbf{x} \cdot \xi^\mu. \quad (2)$$

All these models are energy-based. What sets the LSE energy apart from the classical, polynomial, and exponential models is the logarithm: it acts as a smooth maximum, keeping the energy extensive in N rather than exponentially large, while its gradient is a softmax over patterns, so the retrieval update replaces the state by a softmax-weighted average of the stored patterns, exactly like the attention mechanism of transformers. Model (1) is the Gaussian-kernel member of a family of kernel-based energies. Replacing the Gaussian by the compactly supported Epanechnikov kernel gives the log-sum-ReLU (LSR) model (Hoover et al. 2025).

A natural question is how retrieval survives thermal noise. At a finite temperature T of the thermal bath (distinct from the kernel sharpness β), the boundary $\alpha_c(\beta, T)$ in the load-temperature plane separates cues that relax back to their target from cues that escape to spurious states. Petrova, Polyachenko, and State (2026) derived this boundary for the LSE and LSR kernels under a Gaussian approximation of the inter-pattern alignment density, obtaining for LSE a β -independent zero-temperature load $\alpha_c = 1/2$, and identifying the Gaussian approximation as the likely source of error. Direct Monte Carlo (Petrova 2026) then showed deviations from the Gaussian prediction, with the retrieval region extending to larger α as $T \rightarrow 0$.

Here we revisit the LSE capacity on the N -sphere using the exact spherical alignment density, both analytically and by Monte Carlo, across the full (α, T) plane. The sum over non-target patterns is governed by an interior saddle of the partition-sum integrand rather than by the right edge of the alignment distribution, which restores the β -dependence missing from the Gaussian estimate. At $T = 0$ the resulting capacity reproduces the exact spherical threshold of Lucibello and Mézard (2024) and lies above the rigorous storage lower bound of Ramsauer et al. (2021). At $T > 0$ the retrieval well is bounded by a spinodal, the limit of metastability beyond which the well disappears; below it the well persists metastably, and a trapped cue escapes only over a barrier, with a lifetime exponential in N , a Kramers escape (Kramers 1940; Hänggi, Talkner, and Borkovec 1990). This resolves the low- T deviations seen in simulations.

The paper is organized as follows. Section 2 builds the finite-temperature retrieval theory from the exact spherical alignment density: the Gaussian baseline and why it fails, the interior saddle that fixes the noise floor, the zero-temperature capacity, and the finite-temperature spinodal. The Monte Carlo schemes section describes the direct and truncated samplers. The Results section presents the stability diagram in the (α, T) plane at $\beta = 1$, compares it against the Gaussian and spinodal predictions, and examines the finite- N behavior. The Discussion section collects the implications and the relation to the literature.

Loss of the retrieval in the large- N limit

Let the retrieval pattern ξ^1 and a spurious pattern ξ be drawn independently and uniformly from the sphere of radius $\sqrt{N}$ in $\mathbb{R}^N$. The alignment

$$\varphi(\xi) \equiv N^{-1} \xi^1 \cdot \xi \quad (3)$$

is the cosine of the angle between them. By rotational symmetry we may fix the first pattern along the coordinate axis e_1 and treat the spurious pattern as a unit vector u uniformly distributed on S^{N-1} . The alignment is then simply its first coordinate, $\varphi = u_1$, and we seek the density of u_1 induced by the uniform measure on the sphere. The probability that u_1 lies in $[\varphi, \varphi + d\varphi]$ equals the area of the corresponding spherical band divided by the total area of S^{N-1} . Decomposing $u = (u_1, u_\perp)$, the constraint $|u| = 1$ fixes $|u_\perp| = \sqrt{1 - \varphi^2}$, so the points at height φ form a sphere S^{N-2} of radius $\sqrt{1 - \varphi^2}$, with surface area $\propto (1 - \varphi^2)^{(N-2)/2}$. The band thickness is $d\varphi/\sqrt{1 - \varphi^2}$, since the surface is inclined with respect to the axis. Normalizing the band area $\propto (1 - \varphi^2)^{(N-3)/2} d\varphi$ gives the density

$$h(\varphi) = C_N (1 - \varphi^2)^{(N-3)/2} \quad (4)$$

with the normalization constant

$$C_N = \frac{\Gamma(N/2)}{\sqrt{\pi} \Gamma((N-1)/2)}. \quad (5)$$

The density is symmetric about $\varphi = 0$, with mean zero and variance $1/N$. As $N \rightarrow \infty$ it concentrates near the equator and approaches a Gaussian,

$$h(\varphi) \rightarrow \sqrt{\frac{N}{2\pi}} e^{-N\varphi^2/2}, \quad (6)$$

so distinct high-dimensional patterns are nearly orthogonal: typical alignments are of order $N^{-1/2}$. The exact law is retained where the rare large-alignment tail matters; the Gaussian form is adequate only near the center. The density (4) and its Gaussian limit are the spherical and Gaussian rate functions of Lucibello and Mézard (2024).

Ramsauer et al. (2021) tie retrieval to the separation of a pattern from its nearest neighbour, $N(1 - \varphi_{\max})$, with $\varphi_{\max}$ the highest alignment among the $M = e^{\alpha N}$ patterns (their Theorem 4): retrieval fails when a spurious pattern reaches the target, $\varphi_{\max} \rightarrow 1$. The load at which this happens is fixed by the alignment law $h(\varphi)$ near $\varphi = 1$. Under the Gaussian approximation h keeps weight at $\varphi = 1$, and $\varphi_{\max}$

reaches 1 at the β -independent load $\alpha_c^G = 1/2$. Under the exact law (4) h vanishes at $\varphi = 1$, since two distinct patterns cannot align perfectly on the sphere; then

$$\varphi_{\max} = \sqrt{1 - e^{-2\alpha}} < 1 \quad \text{for every } \alpha, \quad (7)$$

and the capacity is infinite. Both predictions are unsatisfactory: $\alpha_c^G = 1/2$ carries no dependence on the kernel sharpness β , and the exact nearest-pattern estimate gives no finite capacity.

We now analyse retrieval thermodynamically, through the free energy density profile. The free energy density is the internal energy density $u = H/N$ minus T times the entropy. The softmax sum $S(x)$ from (1) splits into a retrieval term and a sum over spurious patterns,

$$S(x) \equiv \exp[\beta N(\varphi(x) - 1)] + S_{\text{spur}}(x), \quad (8)$$

$$S_{\text{spur}}(x) \equiv \sum_{\mu=2}^M \exp[\beta N(\varphi^\mu(x) - 1)]. \quad (9)$$

The spurious patterns are i.i.d. uniform on the sphere and independent of x , so for any x with $|x|^2 = N$ the projections $\varphi^\mu(x) = x \cdot \xi^\mu / N$ have marginal density (4). The spurious sum therefore has the same distribution at every x on the sphere; in particular at $x = \xi^1$, where the projections become the inter-pattern alignments $\varphi_{1\mu}$. With $M = e^{\alpha N}$, the sum concentrates around its expectation, and

$$S_{\text{spur}} \approx M \int_{-1}^1 d\varphi h(\varphi) \exp[\beta N(\varphi - 1)]. \quad (10)$$

The large- N limit of (10) is set by a saddle point. The saddle φ^* solves $\partial g(\varphi, \beta) / \partial \varphi = 0$, with

$$g(\varphi, \beta) = \ln(1 - \varphi^2)/2 + \beta \varphi. \quad (11)$$

This gives the saddle point at

$$\varphi^*(\beta) = \frac{\sqrt{1 + 4\beta^2} - 1}{2\beta}, \quad (12)$$

and the saddle value

$$g_{\max}(\beta) = g(\varphi^*, \beta) = \frac{1}{2} \ln(1 - \varphi^{*2}) + \beta \varphi^*. \quad (13)$$

The leading-order spurious sum is then

$$S_{\text{spur}} \sim \exp[N(\alpha + g_{\max}(\beta) - \beta)]. \quad (14)$$

Appendix A derives the prefactor; it depends only on $\varphi^*(\beta)$ and equals 1.082 at $\beta = 1$. In the thermodynamic limit only the leading term in the exponent (proportional to N) survives, so the prefactor is not needed at this stage. The internal energy density is

$$u(\varphi) = \begin{cases} 1 - \varphi & \text{retrieval term dominates,} \\ 1 - [\alpha + g_{\max}(\beta)]/\beta & \text{spurious term dominates.} \end{cases} \quad (15)$$

They meet in a cusp at $\varphi_c(\alpha, \beta) = (\alpha + g_{\max}(\beta))/\beta$, where dominance switches.

The entropy density (Petrova, Polyachenko, and State 2026)

$$s(\varphi) = \frac{1}{2} \ln(1 - \varphi^2), \quad (16)$$




Figure 1: Leading-order free energy density $f(\varphi)$ at $\alpha = 0.5$, $\beta = 1$, for several T . RB: retrieval basin; SV: saddle valley (where the spurious term dominates); cusp at $\varphi \approx 0.877$ (dashed grey). Coloured dots mark the RB minima; they disappear at the spinodal temperature $T_{\text{sp}}(\alpha, \beta) \approx 0.262$.

gives the free energy density $f(\varphi) = u(\varphi) - T s(\varphi)$, plotted in Fig. 1 for $\alpha = 0.5$, $\beta = 1$ and several T . The profile has two minima: one at $\varphi \approx 0$ from the spurious sum, one at $\varphi \approx 1$ from the retrieval term. The retrieval minimum may be global or local; what matters for retrieval is whether it exists at all, and it disappears at the spinodal. The retrieval minimum is at

$$\varphi_{\text{eq}} = \frac{1}{2}(\sqrt{T^2 + 4} - T), \quad (17)$$

so the spinodal condition $\varphi_{\text{eq}} = \varphi_c(\alpha, \beta)$ yields

$$T_{\text{sp}}(\alpha, \beta) = \frac{\beta^2 - (\alpha + g_{\text{max}}(\beta))^2}{\beta(\alpha + g_{\text{max}}(\beta))}. \quad (18)$$

For $T < T_{\text{sp}}(\alpha, \beta)$ the retrieval minimum exists. It is thermodynamically stable when it is the global minimum of f and metastable when the saddle-valley minimum is deeper. At equilibrium the deeper basin carries the larger weight, with the ratio between the two basins set by their depth difference; at finite T and finite N both basins can carry significant mass. The chain escapes the retrieval basin to the saddle valley by a Kramers process at the same rate $\sim e^{-N\Delta f/T}$ in both regimes (Kramers 1940; Hänggi, Talkner, and Borkovec 1990), with Δf the per-spin barrier from the basin minimum to the cusp; what differs is the return rate. A Monte Carlo ensemble initialised in the basin contains both populations at finite budget, in proportions set by the budget over the Kramers escape time. As $T \rightarrow 0$ the escape rate vanishes, so retrieval is retained whenever the basin exists.

Ramsauer et al. (2021) bound the spurious sum (9) by M times its largest term,

$$S_{\text{spur}} \leq M e^{-\beta N(1-\varphi_{\text{max}})} = e^{N(\alpha - \beta(1-\varphi_{\text{max}}))}, \quad (19)$$

which exceeds the retrieval term once $\varphi_{\text{max}} > 1 - \alpha/\beta$. The highest alignment present is (7), so Ramsauer et al.'s capacity solves

$$\sqrt{1 - e^{-2\alpha}} = 1 - \frac{\alpha}{\beta}, \quad (20)$$




Figure 2: Threshold curves at $\beta = 1$. Red straight line: $\varphi_c(\alpha) = \alpha + g_{\text{max}}(\beta)$, tangent to the blue solid curve at the saddle $\varphi^*(\beta)$ (Appendix B). Blue solid: the exact-density extreme alignment (7). Blue dotted: $\sqrt{2\alpha}$, its Gaussian approximation. Grey dashed: $1 - \alpha/\beta$, the Ramsauer et al. condition.

finite and β -dependent.

Fig. 2 summarises the different approaches to the capacity. The red curve is the zero-temperature capacity from the saddle-point evaluation of the spurious sum. It grows with β and reproduces the exact spherical threshold of Lucibello and Mézard (2024).

The blue solid curve is the exact-density extreme (7) (no finite capacity); the blue dotted curve is its Gaussian approximation $\sqrt{2\alpha}$, which gives the β -independent $\alpha_c^G = 1/2$. The grey dashed line $1 - \alpha/\beta$ is the Ramsauer et al. condition (20); its intersection with the blue solid curve gives the Ramsauer et al. capacity, finite and β -dependent but below the red curve at every β .

The red line is tangent to the blue solid curve at every β ; as β varies, the tangency point sweeps the entire blue solid curve. This follows from the Legendre conjugacy between $g_{\text{max}}(\beta)$ and $\varphi_{\text{max}}(\alpha)$, derived in Appendix B. The geometric consequence is that the cusp $\varphi_c(\alpha, \beta)$ lies above the maximum alignment $\varphi_{\text{max}}(\alpha)$ at every α : no individual spurious pattern enters the retrieval basin $\varphi > \varphi_c$. Retrieval loss is therefore *collective*, driven by the bulk saddle of the spurious sum and not by any single spurious pattern reaching the target. In the terminology of Lucibello and Mézard (2024), the retrieval boundary lies in the *uncondensed* phase of the auxiliary REM, where exponentially many patterns dominate the noise sum, as opposed to the *condensed* phase where a single pattern dominates.

Monte Carlo schemes

Results

Discussion

References

Demircigil, M.; Heusel, J.; Löwe, M.; Uppang, S.; and Vermet, F. 2017. On a Model of Associative Memory with Huge

Storage Capacity. *Journal of Statistical Physics*, 168(2): 288–299.

Gradshteyn, I. S.; and Ryzhik, I. M. 2015. *Table of Integrals, Series, and Products*. Amsterdam: Academic Press, 8 edition. ISBN 978-0-12-384933-5.

Hänggi, P.; Talkner, P.; and Borkovec, M. 1990. Reaction-rate theory: fifty years after Kramers. *Reviews of Modern Physics*, 62(2): 251–341.

Hoover, B.; Shi, Z.; Balasubramanian, K.; Krotov, D.; and Ram, P. 2025. Dense Associative Memory with Epanechnikov Energy. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Hopfield, J. J. 1982. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences*, 79(8): 2554–2558.

Kramers, H. A. 1940. Brownian motion in a field of force and the diffusion model of chemical reactions. *Physica*, 7(4): 284–304.

Krotov, D.; and Hopfield, J. J. 2016. Dense Associative Memory for Pattern Recognition. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 29, 1172–1180.

Lucibello, C.; and Mézard, M. 2024. The Exponential Capacity of Dense Associative Memories. *Physical Review Letters*, 132(7): 077301.

Olver, F. W. J.; Lozier, D. W.; Boisvert, R. F.; and Clark, C. W., eds. 2010. *NIST Handbook of Mathematical Functions*. New York: Cambridge University Press. ISBN 978-0-521-19225-5.

Petrova, T. 2026. Thermal Robustness of Retrieval in Dense Associative Memories: LSE vs LSR Kernels. In *New Frontiers in Associative Memory Workshop at ICLR 2026*. ArXiv:2603.13350.

Petrova, T.; Polyachenko, E.; and State, R. 2026. Geometric Entropy and Retrieval Phase Transitions in Continuous Thermal Dense Associative Memory. In *Proceedings of the 43rd International Conference on Machine Learning (ICML)*, Proceedings of Machine Learning Research. PMLR. ArXiv:2604.07401.

Ramsauer, H.; Schäfl, B.; Lehner, J.; Seidl, P.; Widrich, M.; Gruber, L.; Holzleitner, M.; Adler, T.; Kreil, D.; Kopp, M. K.; Klambauer, G.; Brandstetter, J.; and Hochreiter, S. 2021. Hopfield Networks is All You Need. In *International Conference on Learning Representations*.

Appendix A: Derivation of the prefactor in the spurious sum

The integral (10) evaluates in closed form. Substituting $\varphi = \cos \theta$ maps the spherical density onto the kernel of the integral representation of the modified Bessel function $I_\nu(z)$ (Gradshteyn and Ryzhik 2015, 8.431.3); at $\nu = N/2 - 1$, $z = \beta N$, this gives for $N > 1$

$$S_{\text{spur}} = M\Gamma\left(\frac{N}{2}\right) \left(\frac{2}{\beta N}\right)^{N/2-1} e^{-\beta N} I_{N/2-1}(\beta N). \quad (21)$$

The right-hand side is the disorder expectation of the sum (9). The relative fluctuations of S_{spur} around this expectation are $\sim e^{N[\delta(\beta)-\alpha]/2}$ with $\delta(\beta) = g_{\text{max}}(2\beta) - 2g_{\text{max}}(\beta) > 0$; they are exponentially small in N when $\alpha > \delta(\beta)$. At $\beta = 1$, $\delta \approx 0.334$, well below the retrieval boundary $\alpha_c \approx 0.623$, so the saddle (14) captures the typical S_{spur} .

The Debye uniform asymptotic of the modified Bessel function (Olver et al. 2010, 10.41.3) gives the prefactor of (14). Keeping the leading ($k = 0$) term of the asymptotic series,

$$I_\nu(\nu z') \sim \frac{e^{\nu\eta(z')}}{(2\pi\nu)^{1/2}(1+z'^2)^{1/4}}, \quad (22)$$

where $z' = 2\beta N/(N-2)$ and

$$\eta(z') = \sqrt{1+z'^2} + \ln \frac{z'}{1+\sqrt{1+z'^2}}. \quad (23)$$

$\eta(z')$ is expanded to $\mathcal{O}(N^{-1})$ because it is multiplied by $\nu = N/2 - 1$ in the exponent:

$$\eta(z') = \eta(2\beta) + \eta'(2\beta) \frac{4\beta}{N} + \mathcal{O}(N^{-2}), \quad (24)$$

where

$$\eta(2\beta) = 1 + 2\beta\varphi^* + \ln \varphi^*, \quad (25)$$

$$\eta'(2\beta) = \varphi^* + \frac{1}{2\beta}. \quad (26)$$

Then

$$\nu\eta(z') = \frac{N}{2}\eta(2\beta) - \ln \varphi^* + \mathcal{O}(1/N). \quad (27)$$

The $\mathcal{O}(1)$ piece $-\ln \varphi^*$ in the exponent contributes a factor $1/\varphi^*$ to the prefactor. Using Stirling’s formula

$$\Gamma(N/2) \sim \sqrt{4\pi/N} (N/2)^{N/2} e^{-N/2} \quad (28)$$

and collecting all terms in (21) to $\mathcal{O}(1)$ in the exponent and to leading order in the prefactor reproduces the exponent of (14) and gives the prefactor:

$$S_{\text{spur}} \approx \frac{e^{N[\alpha+g_{\text{max}}(\beta)-\beta]}}{\sqrt{1-\varphi^{*4}}} \quad (29)$$

with $g_{\text{max}}(\beta)$ given by (13). At $\beta = 1$, $\varphi^* \approx 0.618$ and the prefactor is ≈ 1.082 .

Appendix B: Legendre-conjugacy relation between the saddle value and the spherical extreme

The Legendre transform expresses a duality between two ways of describing a convex curve: as the graph of a function $y = f(x)$, or as the envelope of its tangent lines, each labelled by its slope. Legendre introduced it in 1787 to integrate differential equations by changing the independent variable from a coordinate to the corresponding slope; the same construction later became central to thermodynamics (conjugate pairs such as entropy–temperature and volume–pressure) and to convex optimization (primal–dual duality).

Take the simplest convex curve, the parabola $f(x) = x^2/2$. Its tangent at the point $(x_0, x_0^2/2)$ has slope $p = x_0$ and y -intercept $-x_0^2/2 = -p^2/2$. Writing the tangent line as $y = px - p^2/2$ and varying p , the parabola is recovered as the upper envelope of this one-parameter family of lines,

$$\frac{1}{2}x^2 = \sup_p [px - \frac{1}{2}p^2], \quad (30)$$

with the supremum attained at $p = x$. The same construction defines the Legendre transform of an arbitrary convex function $f(x)$,

$$f^*(p) = \sup_x [px - f(x)], \quad (31)$$

and the inverse relation $f(x) = \sup_p [px - f^*(p)]$ recovers f as the upper envelope of its tangent lines indexed by slope. Two convex functions related this way are called *Legendre conjugates*.

The situation in our problem is exactly this. The spherical extreme $\varphi_{\max}(\alpha) = \sqrt{1 - e^{-2\alpha}}$, written inversely as $\alpha = \psi(\varphi)$, takes the form

$$\psi(\varphi) = -\frac{1}{2} \ln(1 - \varphi^2), \quad (32)$$

a convex function on $(-1, 1)$. Its Legendre transform is

$$\psi^*(\beta) = \sup_{\varphi \in (-1, 1)} [\beta \varphi - \psi(\varphi)]. \quad (33)$$

The supremum is attained at the spherical saddle $\varphi = \varphi^*(\beta) = (-1 + \sqrt{1 + 4\beta^2})/(2\beta)$, the unique positive root of

$$\beta \varphi^{*2} + \varphi^* - \beta = 0. \quad (34)$$

Substituting back gives

$$\psi^*(\beta) = \beta \varphi^* + \frac{1}{2} \ln(1 - \varphi^{*2}) = g_{\max}(\beta), \quad (35)$$

the saddle value (13). So $g_{\max}$ and ψ are Legendre conjugates.

The geometric statement of this conjugacy: the cusp line at sharpness β ,

$$\varphi_c(\alpha, \beta) = [\alpha + g_{\max}(\beta)]/\beta, \quad (36)$$

is the tangent line to the spherical extreme with slope $1/\beta$ and α -intercept $-g_{\max}(\beta)$. As β ranges over $(0, \infty)$, the family of cusp lines envelopes the spherical extreme from above,

$$\varphi_{\max}(\alpha) = \inf_{\beta > 0} \varphi_c(\alpha, \beta), \quad (37)$$

with tangency at $\varphi = \varphi^*(\beta)$. The geometry of Fig. 2 is therefore a Legendre duality between the saddle value $g_{\max}$ on the kernel-sharpness side and the spherical extreme $\varphi_{\max}$ on the alignment side.