

Thermodynamic approach and Saddle-Dominated Capacity of the LSE Dense Associative Memory

Anonymous submission

Abstract

The Dense Associative Memory with log-sum-exp (LSE) kernel on the N -sphere stores $M = \exp(\alpha N)$ patterns at load α . A Gaussian approximation of the inter-pattern alignments gives a zero-temperature capacity $\alpha_c = 1/2$, independent of the kernel sharpness β , the Gaussian-ensemble threshold of Lucibello and Mézard (2024). Using the exact spherical alignment density, we build a finite-temperature theory of retrieval. The sum over non-target patterns is then set by an interior saddle rather than by the right tail of the Gaussian, which restores the β -dependence: at $T = 0$ the capacity grows from $\sim \beta$ at small β to $\frac{1}{2} \ln \beta$ at large β , is ≈ 0.6226 at $\beta = 1$, and reproduces the exact spherical threshold of Lucibello and Mézard. At $T > 0$ the retrieval basin is bounded by a single-pattern spinodal, beyond which it no longer exists; below it the basin is metastable, with a Kramers escape time exponential in N . We confirm the resulting (α, T) stability diagram by Monte Carlo: the Full sampler stores all M patterns but reaches only $N \sim 25$, the Cone sampler that keeps only the patterns in a cone around the target reaches $N \sim 30$, and a pattern-free Bulk sampler that follows the leading-order free-energy profile reaches $N \sim 10^4$. All three agree with the predicted spinodal to within the residual finite- N gap.

Introduction

Modern dense associative memory (DAM) models begin with Krotov and Hopfield (2016), who generalized the classical Hopfield network (Hopfield 1982) from pairwise to polynomial interactions, raising the storage capacity from linear to polynomial in the number of neurons N while preserving the ability to retrieve previously stored information from noisy or corrupted input data. Taking the polynomial degree to infinity yields an exponential capacity (Demircigil et al. 2017): for binary neurons (state space $\{-1, +1\}^N$), the limiting capacity is $M_c = \exp(\alpha_c N)$ with load $\alpha_c = (\ln 2)/2 \approx 0.35$.

The same exponential scaling carries over to continuous states. Ramsauer et al. (2021) introduced a modern Hopfield network whose state is a continuous vector in $\mathbb{R}^N$ and whose energy is a log-sum-exp (LSE) of the state-pattern overlaps. We take both the stored patterns ξ^μ and the retrieval state x on the sphere $|x|^2 = N$, and store $M = \exp(\alpha N)$ patterns

at load α . The energy then reads

$$H(x) = -\beta^{-1} \ln \sum_{\mu=1}^M \exp[-\beta N(1 - \phi^\mu)], \quad (1)$$

where β is a kernel sharpness controlling how strongly the energy distinguishes patterns, and ϕ^μ is the *alignment* of the current state x with the μ -th pattern ξ^μ ,

$$\phi^\mu(x) = N^{-1} x \cdot \xi^\mu. \quad (2)$$

All these models are energy-based. What sets the LSE energy apart from the classical, polynomial, and exponential models is the logarithm: it acts as a smooth maximum, keeping the energy extensive in N rather than exponentially large, while its gradient is a softmax over patterns, so the retrieval update replaces the state by a softmax-weighted average of the stored patterns, exactly like the attention mechanism of transformers. Model (1) is the Gaussian-kernel member of a family of kernel-based energies. Replacing the Gaussian by the compactly supported Epanechnikov kernel gives the log-sum-ReLU (LSR) model (Hoover et al. 2025).

A natural question is how retrieval survives thermal noise. At a finite temperature T of the thermal bath (distinct from the kernel sharpness β), the boundary $\alpha_c(\beta, T)$ in the load-temperature plane separates cues that relax back to their target from cues that escape to spurious states. Petrova, Polyachenko, and State (2026) derived this boundary for the LSE and LSR kernels under a Gaussian approximation of the inter-pattern alignment density, obtaining for LSE a β -independent zero-temperature load $\alpha_c = 1/2$, and identifying the Gaussian approximation as the likely source of error. Brute-force Monte Carlo (Petrova 2026) then showed deviations from the Gaussian prediction, with the retrieval region extending to larger α as $T \rightarrow 0$.

Here we revisit the LSE capacity on the N -sphere using the exact spherical alignment density, both analytically and by Monte Carlo, across the full (α, T) plane. The sum over non-target patterns is governed by an interior saddle of the partition-sum integrand rather than by the right edge of the alignment distribution, which restores the β -dependence missing from the Gaussian estimate. At $T = 0$ the resulting capacity reproduces the exact spherical threshold of Lucibello and Mézard (2024) and lies above the rigorous storage lower bound of Ramsauer et al. (2021). At $T > 0$ the

retrieval well is bounded by a spinodal, the limit of metastability beyond which the well disappears; below it the well persists metastably, and a trapped cue escapes only over a barrier, with a lifetime exponential in N , a Kramers escape (Kramers 1940; Hänggi, Talkner, and Borkovec 1990). This resolves the low- T deviations seen in simulations.

The paper is organized as follows. Section builds the finite-temperature retrieval theory from the exact spherical alignment density: the Gaussian baseline and why it fails, the interior saddle that fixes the noise floor, the zero-temperature capacity, and the finite-temperature spinodal. The Monte Carlo schemes section describes the Full, Cone, and Bulk samplers. The Results section presents the stability diagram in the (α, T) plane at $\beta = 1$, compares it against the Gaussian and spinodal predictions, and examines the finite- N behavior. The Discussion section collects the implications and the relation to the literature.

Retrieval at large N

Let the retrieval pattern ξ^1 and a spurious pattern ξ be drawn independently and uniformly from the sphere of radius $\sqrt{N}$ in $\mathbb{R}^N$. The alignment

$$\varphi(\xi) \equiv N^{-1} \xi^1 \cdot \xi \quad (3)$$

is the cosine of the angle between them. By rotational symmetry we may fix the first pattern along the coordinate axis e_1 and treat the spurious pattern as a unit vector $\mathbf{u}$ uniformly distributed on S^{N-1} . The alignment is then simply its first coordinate, $\varphi = u_1$, and we seek the density of u_1 induced by the uniform measure on the sphere. The probability that u_1 lies in $[\varphi, \varphi + d\varphi]$ equals the area of the corresponding spherical band divided by the total area of S^{N-1} . Decomposing $\mathbf{u} = (u_1, \mathbf{u}_\perp)$, the constraint $|\mathbf{u}| = 1$ fixes $|\mathbf{u}_\perp| = \sqrt{1 - \varphi^2}$, so the points at height φ form a sphere S^{N-2} of radius $\sqrt{1 - \varphi^2}$, with surface area $\propto (1 - \varphi^2)^{(N-2)/2}$. The band thickness is $d\varphi/\sqrt{1 - \varphi^2}$, since the surface is inclined with respect to the axis. Normalizing the band area $\propto (1 - \varphi^2)^{(N-3)/2} d\varphi$ gives the density

$$h(\varphi) = C_N (1 - \varphi^2)^{(N-3)/2} \quad (4)$$

with the normalization constant

$$C_N = \frac{\Gamma(N/2)}{\sqrt{\pi} \Gamma((N-1)/2)}. \quad (5)$$

The density is symmetric about $\varphi = 0$, with mean zero and variance $1/N$. As $N \rightarrow \infty$ it concentrates near the equator and approaches a Gaussian,

$$h(\varphi) \longrightarrow \sqrt{\frac{N}{2\pi}} e^{-N\varphi^2/2}, \quad (6)$$

so distinct high-dimensional patterns are nearly orthogonal: typical alignments are of order $N^{-1/2}$. The exact law is retained where the rare large-alignment tail matters; the Gaussian form is adequate only near the center. The density (4) and its Gaussian limit are the spherical and Gaussian rate functions of Lucibello and Mézard (2024).

Extreme-value attempts. Ramsauer et al. (2021) provide two related lower bounds on the storage capacity, both based on extreme-value statistics of the inter-pattern alignments. Their Theorem 4 ties retrieval to the separation of a pattern from its nearest neighbour, $N(1 - \varphi_{\max})$, with $\varphi_{\max}$ the highest alignment among the $M = e^{\alpha N}$ patterns: retrieval fails when a spurious pattern reaches the target, $\varphi_{\max} \rightarrow 1$. The load at which this happens is fixed by the alignment law $h(\varphi)$ near $\varphi = 1$. Under the Gaussian approximation h keeps weight at $\varphi = 1$, and $\varphi_{\max}$ reaches 1 at the β -independent load $\alpha_c^G = 1/2$. Under the exact law (4) h vanishes at $\varphi = 1$, since two distinct patterns cannot align perfectly on the sphere; then

$$\varphi_{\max} = \sqrt{1 - e^{-2\alpha}} < 1 \quad \text{for every } \alpha, \quad (7)$$

and the capacity is infinite.

Their Theorem 3 acts directly on the softmax sum $S(\mathbf{x})$ from (1), which splits into a retrieval term and a sum over spurious patterns,

$$S(\mathbf{x}) \equiv \exp[\beta N(\varphi(\mathbf{x}) - 1)] + S_{\text{spur}}(\mathbf{x}), \quad (8)$$

$$S_{\text{spur}}(\mathbf{x}) \equiv \sum_{\mu=2}^M \exp[\beta N(\varphi^\mu(\mathbf{x}) - 1)]. \quad (9)$$

The bound replaces S_{spur} by M times its largest term,

$$S_{\text{spur}} \leq M e^{-\beta N(1 - \varphi_{\max})} = e^{N(\alpha - \beta(1 - \varphi_{\max}))}, \quad (10)$$

which exceeds the retrieval term once $\varphi_{\max} > 1 - \alpha/\beta$. With $\varphi_{\max}$ from (7), Ramsauer et al.'s capacity solves

$$\sqrt{1 - e^{-2\alpha}} = 1 - \frac{\alpha}{\beta}, \quad (11)$$

finite and β -dependent.

Neither bound is tight. Theorem 4 with the Gaussian gives a β -independent $1/2$; with the exact density it gives infinity. Theorem 3 gives a finite β -dependent capacity but assumes S_{spur} is dominated by its largest term. The saddle-point analysis below evaluates S_{spur} without that approximation.

Saddle-point evaluation. The spurious patterns are i.i.d. uniform on the sphere and independent of $\mathbf{x}$, so for any $\mathbf{x}$ with $|\mathbf{x}|^2 = N$ the projections $\varphi^\mu(\mathbf{x}) = \mathbf{x} \cdot \xi^\mu/N$ have marginal density (4). The spurious sum therefore has the same distribution at every $\mathbf{x}$ on the sphere; in particular at $\mathbf{x} = \xi^1$, where the projections become the inter-pattern alignments $\varphi_{1\mu}$. With $M = e^{\alpha N}$, the sum concentrates around its expectation, and

$$S_{\text{spur}} \approx M \int_{-1}^1 d\varphi h(\varphi) \exp[\beta N(\varphi - 1)]. \quad (12)$$

The large- N limit of (12) is set by a saddle point. The saddle φ^* solves $\partial g(\varphi, \beta)/\partial \varphi = 0$, with

$$g(\varphi, \beta) = \ln(1 - \varphi^2)/2 + \beta \varphi. \quad (13)$$

This gives the saddle point at

$$\varphi^*(\beta) = \frac{\sqrt{1 + 4\beta^2} - 1}{2\beta}, \quad (14)$$

and the saddle value

$$g_{\max}(\beta) = g(\varphi^*, \beta) = \frac{1}{2} \ln(1 - \varphi^{*2}) + \beta \varphi^*. \quad (15)$$

The leading-order spurious sum is then

$$S_{\text{spur}} \sim \exp[N(\alpha + g_{\max}(\beta) - \beta)]. \quad (16)$$

Appendix A derives the prefactor; it depends only on $\varphi^*(\beta)$ and equals 1.082 at $\beta = 1$. In the large- N limit only the leading term in the exponent (proportional to N) survives, so the prefactor is not needed at this stage.

Zero-temperature capacity. The internal energy density $u = H/N$ from (1) reads

$$u(\varphi) = \begin{cases} 1 - \varphi & \text{retrieval term dominates,} \\ 1 - [\alpha + g_{\max}(\beta)]/\beta & \text{spurious term dominates.} \end{cases} \quad (17)$$

The two branches meet in a cusp at

$$\varphi_c(\alpha, \beta) = (\alpha + g_{\max}(\beta))/\beta, \quad (18)$$

where dominance switches. At zero temperature the retrieval minimum sits at $\varphi = 1$; it is the global minimum as long as $\varphi_c < 1$, giving the zero-temperature capacity

$$\alpha_c^{\text{sd}}(\beta, 0) = \beta - g_{\max}(\beta). \quad (19)$$

This boundary matches the exact spherical threshold of Lucibello and Mézard (2024) at every β .

Geometric picture. Fig. 1 summarises the different approaches at $\beta = 1$. The red straight line is the cusp position $\varphi_c(\alpha)$; the saddle-point capacity (19) is the load at which it reaches $\varphi = 1$. The blue solid curve is the exact-density extreme (7) (no finite capacity); the blue dotted curve is its Gaussian approximation $\sqrt{2\alpha}$, which gives the β -independent $\alpha_c^G = 1/2$. The grey dashed line $1 - \alpha/\beta$ is the Ramsauer et al. condition (11); its intersection with the blue solid curve gives the Ramsauer et al. capacity, finite and β -dependent but below the saddle-point capacity at every β .

The red line is tangent to the blue solid curve at every β ; as β varies, the tangency point sweeps the entire blue solid curve. This follows from the Legendre conjugacy between $g_{\max}(\beta)$ and $\varphi_{\max}(\alpha)$, derived in Appendix B. The geometric consequence is that the cusp $\varphi_c(\alpha, \beta)$ lies above the maximum alignment $\varphi_{\max}(\alpha)$ at every α : no individual spurious pattern enters the retrieval basin $\varphi > \varphi_c$. Retrieval loss is therefore *collective*, driven by the bulk saddle of the spurious sum and not by any single spurious pattern reaching the target. In the terminology of Lucibello and Mézard (2024), the retrieval boundary lies in the *uncondensed* phase of the auxiliary REM, where exponentially many patterns dominate the noise sum, as opposed to the *condensed* phase where a single pattern dominates.

Extension to finite temperature. At $T > 0$ the entropy density (Petrova, Polyachenko, and State 2026)

$$s(\varphi) = \frac{1}{2} \ln(1 - \varphi^2), \quad (20)$$

adds to the internal energy, giving the free energy density $f(\varphi) = u(\varphi) - T s(\varphi)$, plotted in Fig. 2 for $\alpha = 0.5$, $\beta = 1$




Figure 1: Threshold curves at $\beta = 1$. Red straight line: $\varphi_c(\alpha) = \alpha + g_{\max}(\beta)$, tangent to the blue solid curve at the saddle $\varphi^*(\beta)$ (Appendix B). Blue solid: the exact-density extreme alignment (7). Blue dotted: $\sqrt{2\alpha}$, its Gaussian approximation. Grey dashed: $1 - \alpha/\beta$, the Ramsauer et al. condition.

and several T . The profile has two minima: one at $\varphi \approx 0$ from the spurious sum, one at $\varphi \approx 1$ from the retrieval term. The retrieval minimum shifts from $\varphi = 1$ at $T = 0$ to

$$\varphi_{\text{eq}}(T) = \frac{1}{2}(\sqrt{T^2 + 4} - T) \quad (21)$$

at $T > 0$, and disappears at the spinodal $\varphi_{\text{eq}}(T) = \varphi_c(\alpha, \beta)$, giving

$$T_{\text{sp}}(\alpha, \beta) = \frac{\beta^2 - (\alpha + g_{\max}(\beta))^2}{\beta(\alpha + g_{\max}(\beta))}. \quad (22)$$

At fixed T , the saddle-dominated capacity is the load at which the retrieval free energy $f_{\text{ret}}(T) = (1 - \varphi_{\text{eq}}) - \frac{T}{2} \ln(1 - \varphi_{\text{eq}}^2)$ equals the saddle-valley value, giving

$$\alpha_c^{\text{sd}}(\beta, T) = \beta[1 - f_{\text{ret}}(T)] - g_{\max}(\beta), \quad (23)$$

which reduces to (19) at $T = 0$.

Stability and Kramers escape. For $T < T_{\text{sp}}(\alpha, \beta)$ the retrieval minimum exists. It is thermodynamically stable when it is the global minimum of f and metastable when the saddle-valley minimum is deeper. At equilibrium the deeper basin carries the larger weight, with the ratio between the two basins set by their depth difference; at finite T and finite N both basins can carry significant mass. The chain escapes the retrieval basin to the saddle valley by a Kramers process at the same rate $\sim e^{-N\Delta f/T}$ in both regimes (Kramers 1940; Hänggi, Talkner, and Borkovec 1990), with Δf the per-spin barrier from the basin minimum to the cusp; what differs is the return rate. A Monte Carlo ensemble initialised in the basin contains both populations at finite budget, in proportions set by the budget over the Kramers escape time. As $T \rightarrow 0$ the escape rate vanishes, so retrieval is retained whenever the basin exists.

Monte Carlo schemes

We compare three Metropolis schemes on $\mathbf{x} \in \sqrt{N} \cdot S^{N-1}$ at $\beta = 1$. Each step proposes $\mathbf{x} + \sigma \boldsymbol{\eta}$ with $\boldsymbol{\eta} \sim \mathcal{N}(0, I_N)$




Figure 2: Leading-order free energy density $f(\varphi)$ at $\alpha = 0.5$, $\beta = 1$, for several T . RB: retrieval basin; SV: saddle valley (where the spurious term dominates); cusp at $\varphi \approx 0.877$ (dashed grey). Coloured dots mark the RB minima; they disappear at the spinodal temperature $T_{\text{sp}}(\alpha, \beta) \approx 0.262$.

and rescales the result to $|\mathbf{x}| = \sqrt{N}$. The (α, T) grid in this paper is $\alpha \in \{0.20, 0.21, \dots, 0.65\}$ (46 loads) and $T \in \{0.005, 0.015, \dots, 0.495\}$ (50 temperatures). The schemes differ in how they represent the $M = e^{\alpha N}$ patterns inside the LSE log-sum.

Full. All $M = \lceil e^{\alpha N} \rceil$ patterns are generated explicitly per disorder realization and the log-sum is evaluated over all of them at every step. Cost per step is $\mathcal{O}(MN)$, dominated by the GEMM that produces the M pattern-state overlaps. We fix the budget at $M = 4.4 \times 10^6$ and slide N with α along the ramp $N(\alpha) = \lceil \log M / \alpha \rceil$ (floor $N \geq 12$), so the Full sweep matches the Cone sweep cell-by-cell. The ramp runs from $N = 76$ at $\alpha = 0.20$ to $N = 24$ at $\alpha = 0.65$. At each (α, T) we draw 32 disorder realizations and follow two Metropolis replicas per realization, both initialized at the target ξ^1 . Step size is $\sigma = 2.4 T / \sqrt{N}$. Each replica is equilibrated for $N_{\text{eq}} = 2^{15}$ steps and sampled for $N_{\text{samp}} = 2^{13}$ steps. Per cell we record the time-averaged target overlaps φ_a, φ_b , the replica overlap q_{ab} , and the maximum non-target overlap. The full sweep takes ≈ 26 days on a single NVIDIA RTX A6000 at Float32.

Cone. A fixed alignment cutoff ϕ_{keep} partitions the pattern set into two pieces. Patterns with $\varphi_{1\mu} \geq \phi_{\text{keep}}$ lie inside a spherical cone of half-angle $\arccos(\phi_{\text{keep}})$ around the target ξ^1 ; these are generated explicitly and enter the log-sum term-by-term. Patterns with $\varphi_{1\mu} < \phi_{\text{keep}}$ make up the passive bulk and are absorbed into a single $\mathbf{x}$ -independent constant

$$\ln C_{\text{bulk}} = \alpha N + \ln \int_{-1}^{\phi_{\text{keep}}} h(\phi) e^{\beta N \phi} d\phi, \quad (24)$$

which sits inside the log-sum on equal footing with the in-cone patterns. The number of in-cone patterns per disorder is a Poisson draw with mean $M p_{>}$, where $p_{>} = \Pr(\varphi \geq \phi_{\text{keep}})$ under the spherical density (??). In the rescaled variable $y = (1 + \phi)/2$, the spherical density becomes a symmetric Beta and each in-cone pattern's overlap with ξ^1 is

drawn from the truncated distribution

$$y \sim \text{Beta}\left(\frac{N-1}{2}, \frac{N-1}{2}\right) \text{ truncated to } y \in [(1+\phi_{\text{keep}})/2, 1], \quad \varphi_{1\mu} = \quad (25)$$

together with a uniform perpendicular direction on the $(N-2)$ -sphere orthogonal to ξ^1 . The cone set is fixed once per disorder. We use $M = 3 \times 10^8$ and $\phi_{\text{keep}} = 0.40$ with the same $N(\alpha)$ ramp formula as Full; the ramp runs from $N = 98$ at $\alpha = 0.20$ to $N = 30$ at $\alpha = 0.65$, and the typical in-cone count grows from $K \sim 7 \times 10^3$ at $\alpha = 0.20$ to $K \sim 4 \times 10^6$ at $\alpha = 0.65$, i.e. about 2%–5% of M typically and never more than 5% at any cell. Restricting the explicit pattern set to this small cone is what makes the scheme tractable at $M = 3 \times 10^8$; the remaining $\sim 98\%$ of the patterns enter through the closed-form bulk integral. Disorder count, replicas, initial condition, step size, equilibration and sampling lengths, and recorded observables are identical to Full. At the shared budget $M = 4.4 \times 10^6$, $N = 76$, Full and Cone boundaries agree across the α range we measure, so the bulk patterns are genuinely outside the basin physics rather than silently dropped. The full Cone sweep at $M = 3 \times 10^8$ takes ≈ 7 days on the same A6000.

Bulk. No patterns are generated. Cone, taken in the limit $\phi_{\text{keep}} \rightarrow 1$, admits no patterns inside the cone at all; the integral in the bulk constant (24) simultaneously covers the entire spherical density, and in the large- N limit it is dominated by its saddle point, giving the closed form $M e^{N g_{\text{max}}(\beta)}$. We name this scheme *Bulk* because the log-sum it leaves the chain with consists of exactly two terms, the target and the saddle-evaluated bulk constant: no pattern survives explicitly. The chain energy depends on $\mathbf{x}$ only through the target overlap $\varphi_1 = \xi^1 \cdot \mathbf{x} / N$:

$$\ln Z(\varphi_1) = \ln[e^{\beta N \varphi_1} + e^{N(\alpha + g_{\text{max}}(\beta))}]. \quad (26)$$

Cost per step is $\mathcal{O}(N)$, set by the proposal and the sphere projection, with no M -dependence at all. The chain starts inside the basin at $\varphi_1 = \phi_{\text{eq}}(T)$ with perpendicular components uniform on the $(N-1)$ -sphere of radius $\sqrt{N} \cdot \sqrt{1 - \phi_{\text{eq}}^2}$; an init at $\varphi_1 = 1$ stalls because σ sized for basin diffusion is too aggressive for the linear slope at the sphere edge. The step size $\sigma = \sqrt{2T/\gamma_{\text{eq}}}$ with $\gamma_{\text{eq}} = T(1 + \phi_{\text{eq}}^2)/(1 - \phi_{\text{eq}}^2)^2$ is set from the basin curvature so that the per-step energy change is $\mathcal{O}(T)$ at any N , giving N -independent acceptance. There is no disorder; trial randomness comes only from the MC RNG. Per α we scan T downward from $\min(T_{\text{sp}}(\alpha), 0.495)$ in steps of 0.01, running 64 independent chains for $T_{\text{MC}} = 10^5$ steps each, and record the boundary $T_{\text{bd}}(\alpha)$ as the largest T at which the residual $\phi_{\text{eq}}(T) - \langle \varphi_1 \rangle$ first stays below 0.02 for three consecutive T -steps. We run Bulk at $N \in \{100, 300, 10000\}$ to expose the kinetic boundary as a function of dimension; on a single CPU node with 16 threads the three sweeps take about a minute, about five minutes, and about ninety minutes respectively.

N -reach and what each scheme tests. Cone at fixed ϕ_{keep} keeps $K = M \Pr(\phi \geq \phi_{\text{keep}}) \sim e^{N[\alpha - I(\phi_{\text{keep}})]}$ pat-

terns inside the cone, with $I(\phi) \equiv -\frac{1}{2} \ln(1 - \phi^2)$ the large- N rate of the spherical tail. At $\phi_{\text{keep}} = 0.4$, $I \approx 0.087$; for any $\alpha > 0.087$ the in-cone count grows exponentially in N and the GPU budget is exhausted around $N \sim 30\text{--}50$ at the $M = 3 \times 10^8$ Cone budget. Bulk sets the cone aperture to zero ($\phi_{\text{keep}} \rightarrow 1$) and the bulk integral to its saddle value $M e^{N g_{\text{max}}(\beta)}$, so $K = 0$ by construction; the chain energy reduces to the leading-order log-sum (26) and runs at $\mathcal{O}(N)$ cost per step, decoupling from M entirely.

The three schemes probe different physics. Bulk realises the Kramers kinetics of (??) at large N , but the analytic bulk constant $M e^{N g_{\text{max}}(\beta)}$ is taken on faith and any correction to g_{max} would be invisible. Cone with ϕ_{keep} near or below the integrand saddle $\phi^*(\beta)$ retains the patterns that dominate the spurious-sum integral and audits the bulk directly; a subleading correction to g_{max} would appear as a mismatch between Cone and Full boundaries. At the shared budget $M = 4.4 \times 10^6$ the two agree across our α range, so the leading-order analytic bulk captures the spurious physics to within MC noise at the boundary.

Why local refresh fails. A faster scheme (sometimes called “smart”) refreshes the in-cone set every N_{refresh} steps from a cone around the current chain position, keeping only the patterns whose alignment with $\mathbf{x}$ is currently large. The refresh breaks the quenched-disorder assumption that underlies (??)–(??). As $\mathbf{x}$ wanders, new patterns are minted at every refresh, so the cumulative set of patterns seen along a single trajectory exceeds M , and the chain effectively faces a denser pattern set than the model defines. In the retrieval basin the spurious-recruitment rate alone destroys retrieval, for reasons unrelated to (??). Full, Cone, and Bulk all avoid the issue: Full uses every pattern at every step, Cone fixes the in-cone set once per disorder, and Bulk carries no patterns at all.

Results

Fig. 3 shows the empirical (α, T) stability diagram at $\beta = 1$, from the Cone Monte Carlo sampler at two budgets, $M = 4.4 \times 10^6$ and $M = 3.0 \times 10^8$. The heatmap encodes the residual $\varphi_{\text{eq}}(T) - \langle \varphi \rangle$ from the $M = 3.0 \times 10^8$ run: white where the chain remains in the retrieval basin, red where it escapes to the saddle valley. The blue staircases trace, for each α column, the lowest T at which roughly 5% of chain endpoints have escaped from the retrieval basin to the saddle valley. Operationally, this is the threshold $\varphi_{\text{eq}}(T) - \langle \varphi \rangle > 0.02$ at the empirical basin-to-saddle gap $\Delta\varphi \approx 0.40$ along the boundary, after a light smoothing pass to suppress single-cell MC noise. At fixed α , increasing the budget from $M = 4.4 \times 10^6$ to $M = 3.0 \times 10^8$ (which raises the reachable N via the N -ramp) shifts the empirical boundary to higher T . The expected destination at $N \rightarrow \infty$ is the spinodal $T_{\text{sp}}(\alpha, \beta)$ (22): below the spinodal the retrieval basin still exists, and a chain initialised in the basin escapes only on the Kramers time $\tau_K \sim e^{N\Delta f/T}$, which grows exponentially with N . At the N values reached by the Cone sampler here ($N \leq 98$) the boundary is still far from the spinodal. It happens to lie close to the saddle line $\alpha_c^{\text{sd}}(\beta, T)$ (23), but this proximity is incidental: the saddle

line is the thermodynamic crossover where the basin and the saddle valley have equal free energy, and chains in an MC ensemble redistribute continuously between the two minima as α crosses it, with no sharp change in the retrieval fraction. The Gaussian prediction $\alpha_c^G(T)$ sits to the left of both Cone boundaries, underestimating the capacity at every T .

The green curves trace the Bulk boundary at the same threshold, $T_{\text{MC}} = 10^5$ steps per chain, and $N \in \{100, 300, 10^4\}$. Bulk carries no live patterns ($K = 0$) and the chain follows the two-term log-sum (26) directly, so the line is the pure Kramers-escape boundary from the basin minimum at finite N . Empirically the gap to the spinodal shrinks as $1/\sqrt{N}$: at $\alpha = 0.50$ the gap drops from $\Delta T \approx 0.13$ at $N = 100$ to 0.08 at $N = 300$, and by $N = 10^4$ the boundary sits within one T -grid cell of the spinodal across $\alpha \in [0.40, 0.62]$ and is visually indistinguishable from it (the spinodal in the figure is drawn last, on top of the lightest green $N = 10^4$ curve, by design). The march of the Bulk boundary onto the spinodal as N grows is the asymptotic claim the theory makes, now visible across two decades of N in a single figure.

Dependence of the Bulk boundary on the MC budget is much weaker than on N . Solving the 5% criterion $T_{\text{MC}} k = 0.05$ for the Kramers rate $k = (1/2\pi) e^{-N\Delta f/T}$ gives $N\Delta f/T = \ln T_{\text{MC}} + \mathcal{O}(1)$, so the gap to the spinodal grows only as $\sqrt{\ln T_{\text{MC}}}$: about 9% per decade in T_{MC} from the present $T_{\text{MC}} = 10^5$, against the $\sqrt{10} \approx 3.16$ shrinkage per decade in N . Concretely, going from $T_{\text{MC}} = 10^5$ to 10^6 would shift the $N = 100$ Bulk line by $\Delta T_{\text{bd}} \approx 0.01\text{--}0.04$ at $T \geq 0.2$ and by less than the T -grid spacing $dT = 0.01$ at $T \leq 0.05$. Near the LSE capacity tip $\alpha_{\text{LSE}} \approx 0.6226$ the spinodal itself collapses to $T_{\text{sp}} = 0$, so the absolute gap budget is small and T_{MC} shifts there are pinned below the T -grid spacing. Convergence to the spinodal in the figure is therefore controlled by N ; multiplying the per-chain budget by orders of magnitude would barely move any of the three green curves.

Discussion

References

- Demircigil, M.; Heusel, J.; Löwe, M.; Uppang, S.; and Vermet, F. 2017. On a Model of Associative Memory with Huge Storage Capacity. *Journal of Statistical Physics*, 168(2): 288–299.
- Gradshteyn, I. S.; and Ryzhik, I. M. 2015. *Table of Integrals, Series, and Products*. Amsterdam: Academic Press, 8 edition. ISBN 978-0-12-384933-5.
- Hänggi, P.; Talkner, P.; and Borkovec, M. 1990. Reaction-rate theory: fifty years after Kramers. *Reviews of Modern Physics*, 62(2): 251–341.
- Hoover, B.; Shi, Z.; Balasubramanian, K.; Krotov, D.; and Ram, P. 2025. Dense Associative Memory with Epanechnikov Energy. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Hopfield, J. J. 1982. Neural networks and physical systems with emergent collective computational abilities. *Pro-*




Figure 3: Empirical (α, T) stability diagram at $\beta = 1$. Heatmap: $\varphi_{\text{eq}}(T) - \langle \varphi \rangle$ from the $M = 3.0 \times 10^8$ Cone MC run. Blue staircases: Cone retrieval boundary at threshold 0.02 (roughly 5% of chain endpoints escaped) for $M = 4.4 \times 10^6$ (dark) and $M = 3.0 \times 10^8$ (light). Green curves: Bulk retrieval boundary at $T_{\text{MC}} = 10^5$ and $N \in \{100, 300, 10^4\}$, dark to light with N . The light-green $N = 10^4$ curve sits on the spinodal across most of the α range; the spinodal is drawn last to make this overlap visible. Reference curves: saddle boundary $\alpha_c^{\text{sd}}(\beta, T)$ (23) (red dotted), spinodal T_{sp} (22) (black dotted), Gaussian baseline $\alpha_c^{\text{G}}(T)$ (grey dash-dotted).

ceedings of the National Academy of Sciences, 79(8): 2554–2558.

Kramers, H. A. 1940. Brownian motion in a field of force and the diffusion model of chemical reactions. *Physica*, 7(4): 284–304.

Krotov, D.; and Hopfield, J. J. 2016. Dense Associative Memory for Pattern Recognition. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 29, 1172–1180.

Lucibello, C.; and Mézard, M. 2024. The Exponential Capacity of Dense Associative Memories. *Physical Review Letters*, 132(7): 077301.

Olver, F. W. J.; Lozier, D. W.; Boisvert, R. F.; and Clark, C. W., eds. 2010. *NIST Handbook of Mathematical Functions*. New York: Cambridge University Press. ISBN 978-0-521-19225-5.

Petrova, T. 2026. Thermal Robustness of Retrieval in Dense Associative Memories: LSE vs LSR Kernels. In *New Frontiers in Associative Memory Workshop at ICLR 2026*. ArXiv:2603.13350.

Petrova, T.; Polyachenko, E.; and State, R. 2026. Geometric Entropy and Retrieval Phase Transitions in Continuous Thermal Dense Associative Memory. In *Proceedings of the 43rd International Conference on Machine Learning (ICML)*, Proceedings of Machine Learning Research. PMLR. ArXiv:2604.07401.

Ramsauer, H.; Schäfl, B.; Lehner, J.; Seidl, P.; Widrich, M.; Gruber, L.; Holzleitner, M.; Adler, T.; Kreil, D.; Kopp, M. K.; Klambauer, G.; Brandstetter, J.; and Hochreiter, S. 2021. Hopfield Networks is All You Need. In *International Conference on Learning Representations*.

Appendix A: Derivation of the prefactor in the spurious sum

The integral (12) evaluates in closed form. Substituting $\varphi = \cos \theta$ maps the spherical density onto the kernel of the integral representation of the modified Bessel function $I_\nu(z)$ (Gradshteyn and Ryzhik 2015, 8.431.3); at $\nu = N/2 - 1$, $z = \beta N$, this gives for $N > 1$

$$S_{\text{spur}} = MT \left(\frac{N}{2} \right) \left(\frac{2}{\beta N} \right)^{N/2-1} e^{-\beta N} I_{N/2-1}(\beta N). \quad (27)$$

The right-hand side is the disorder expectation of the sum (9). The relative fluctuations of S_{spur} around this expectation are $\sim e^{N[\delta(\beta) - \alpha]/2}$ with $\delta(\beta) = g_{\text{max}}(2\beta) - 2g_{\text{max}}(\beta) > 0$; they are exponentially small in N when $\alpha > \delta(\beta)$. At $\beta = 1$, $\delta \approx 0.334$, well below the retrieval boundary $\alpha_c \approx 0.623$, so the saddle (16) captures the typical S_{spur} .

The Debye uniform asymptotic of the modified Bessel function (Olver et al. 2010, 10.41.3) gives the prefactor of (16). Keeping the leading ($k = 0$) term of the asymptotic series,

$$I_\nu(\nu z') \sim \frac{e^{\nu \eta(z')}}{(2\pi\nu)^{1/2}(1+z'^2)^{1/4}}, \quad (28)$$

where $z' = 2\beta N/(N - 2)$ and

$$\eta(z') = \sqrt{1+z'^2} + \ln \frac{z'}{1+\sqrt{1+z'^2}}. \quad (29)$$

$\eta(z')$ is expanded to $\mathcal{O}(N^{-1})$ because it is multiplied by $\nu = N/2 - 1$ in the exponent:

$$\eta(z') = \eta(2\beta) + \eta'(2\beta) \frac{4\beta}{N} + \mathcal{O}(N^{-2}), \quad (30)$$

where

$$\eta(2\beta) = 1 + 2\beta\varphi^* + \ln \varphi^*, \quad (31)$$

$$\eta'(2\beta) = \varphi^* + \frac{1}{2\beta}. \quad (32)$$

Then

$$\nu\eta(z') = \frac{N}{2}\eta(2\beta) - \ln \varphi^* + \mathcal{O}(1/N). \quad (33)$$

The $\mathcal{O}(1)$ piece $-\ln \varphi^*$ in the exponent contributes a factor $1/\varphi^*$ to the prefactor. Using Stirling's formula

$$\Gamma(N/2) \sim \sqrt{4\pi/N} (N/2)^{N/2} e^{-N/2} \quad (34)$$

and collecting all terms in (27) to $\mathcal{O}(1)$ in the exponent and to leading order in the prefactor reproduces the exponent of (16) and gives the prefactor:

$$S_{\text{spur}} \approx \frac{e^{N[\alpha + g_{\max}(\beta) - \beta]}}{\sqrt{1 - \varphi^{*4}}} \quad (35)$$

with $g_{\max}(\beta)$ given by (15). At $\beta = 1$, $\varphi^* \approx 0.618$ and the prefactor is ≈ 1.082 .

Appendix B: Legendre-conjugacy relation between the saddle value and the spherical extreme

The Legendre transform expresses a duality between two ways of describing a convex curve: as the graph of a function $y = f(x)$, or as the envelope of its tangent lines, each labelled by its slope. Legendre introduced it in 1787 to integrate differential equations by changing the independent variable from a coordinate to the corresponding slope; the same construction later became central to thermodynamics (conjugate pairs such as entropy–temperature and volume–pressure) and to convex optimization (primal–dual duality).

Take the simplest convex curve, the parabola $f(x) = x^2/2$. Its tangent at the point $(x_0, x_0^2/2)$ has slope $p = x_0$ and y -intercept $-x_0^2/2 = -p^2/2$. Writing the tangent line as $y = px - p^2/2$ and varying p , the parabola is recovered as the upper envelope of this one-parameter family of lines,

$$\frac{1}{2}x^2 = \sup_p [px - \frac{1}{2}p^2], \quad (36)$$

with the supremum attained at $p = x$. The same construction defines the Legendre transform of an arbitrary convex function $f(x)$,

$$f^*(p) = \sup_x [px - f(x)], \quad (37)$$

and the inverse relation $f(x) = \sup_p [px - f^*(p)]$ recovers f as the upper envelope of its tangent lines indexed by slope. Two convex functions related this way are called *Legendre conjugates*.

The situation in our problem is exactly this. The spherical extreme $\varphi_{\max}(\alpha) = \sqrt{1 - e^{-2\alpha}}$, written inversely as $\alpha = \psi(\varphi)$, takes the form

$$\psi(\varphi) = -\frac{1}{2} \ln(1 - \varphi^2), \quad (38)$$

a convex function on $(-1, 1)$. Its Legendre transform is

$$\psi^*(\beta) = \sup_{\varphi \in (-1, 1)} [\beta\varphi - \psi(\varphi)]. \quad (39)$$

The supremum is attained at the spherical saddle $\varphi = \varphi^*(\beta) = (-1 + \sqrt{1 + 4\beta^2})/(2\beta)$, the unique positive root of

$$\beta\varphi^{*2} + \varphi^* - \beta = 0. \quad (40)$$

Substituting back gives

$$\psi^*(\beta) = \beta\varphi^* + \frac{1}{2} \ln(1 - \varphi^{*2}) = g_{\max}(\beta), \quad (41)$$

the saddle value (15). So $g_{\max}$ and ψ are Legendre conjugates.

The geometric statement of this conjugacy: the cusp line at sharpness β ,

$$\varphi_c(\alpha, \beta) = [\alpha + g_{\max}(\beta)]/\beta, \quad (42)$$

is the tangent line to the spherical extreme with slope $1/\beta$ and α -intercept $-g_{\max}(\beta)$. As β ranges over $(0, \infty)$, the family of cusp lines envelopes the spherical extreme from above,

$$\varphi_{\max}(\alpha) = \inf_{\beta > 0} \varphi_c(\alpha, \beta), \quad (43)$$

with tangency at $\varphi = \varphi^*(\beta)$. The geometry of Fig. 1 is therefore a Legendre duality between the saddle value $g_{\max}$ on the kernel-sharpness side and the spherical extreme $\varphi_{\max}$ on the alignment side.

Appendix C: Tangency of the cusp lines to the spherical extreme, in figure coordinates

This appendix shows that the cusp lines $\varphi_c(\alpha, \beta) = (\alpha + g_{\max}(\beta))/\beta$ are tangent to the spherical-extreme curve $\varphi_{\max}(\alpha) = \sqrt{1 - e^{-2\alpha}}$ at every β , that the tangency point sweeps the full curve as β ranges over $(0, \infty)$, and that the spherical extreme is the lower envelope of the cusp-line family. The argument works directly in the (α, φ) coordinates of Fig. 1: α on the x-axis, φ on the y-axis. It splits into two steps: matching the slope of the tangent to the blue curve with the slope of the cusp line, which gives parallelism, and matching the y -intercept, which upgrades parallelism to coincidence. The slope step pulls out the saddle equation; the intercept step expresses the Legendre conjugacy of $g_{\max}$ and $\varphi_{\max}$.

Concavity of the blue curve. Differentiating $\varphi_{\max}(\alpha) = \sqrt{1 - e^{-2\alpha}}$ once gives

$$\varphi'_{\max}(\alpha) = \frac{1 - \varphi_{\max}^2}{\varphi_{\max}}, \quad (44)$$

and a second differentiation gives $\varphi''_{\max}(\alpha) < 0$ for $\alpha > 0$, so $\varphi_{\max}$ is strictly concave on $(0, \infty)$. For a concave function every tangent line lies above the function, with equality only at the tangency point. This will give the lower-envelope identity at the end.

Step A: matching the slope (parallelism). The cusp line at sharpness β ,

$$\varphi = \alpha/\beta + g_{\max}(\beta)/\beta, \quad (45)$$

has slope $1/\beta$ and y-intercept $g_{\max}(\beta)/\beta$. To find the point on the blue curve where the tangent has the same slope, set the right-hand side of (44) equal to $1/\beta$:

$$\frac{1 - \varphi^2}{\varphi} = \frac{1}{\beta} \iff \beta \varphi^2 + \varphi - \beta = 0. \quad (46)$$

This is the saddle equation; its unique positive root is the spherical saddle

$$\varphi^*(\beta) = \frac{-1 + \sqrt{1 + 4\beta^2}}{2\beta}. \quad (47)$$

The corresponding α on the blue curve follows from $\varphi^* = \varphi_{\max}(\alpha^*)$, that is, $e^{-2\alpha^*} = 1 - \varphi^{*2}$:

$$\alpha^*(\beta) = -\frac{1}{2} \ln(1 - \varphi^{*2}). \quad (48)$$

At the point $(\alpha^*(\beta), \varphi^*(\beta))$ on the blue curve, the tangent has slope $1/\beta$, equal to the slope of the cusp line. Step A establishes *parallelism* between the tangent at this point and the cusp line; it does not yet show the two lines coincide.

Step B: matching the y-intercept (coincidence). A line of slope $1/\beta$ passing through $(\alpha^*(\beta), \varphi^*(\beta))$ has y-intercept

$$b_{\text{tangent}}(\beta) = \varphi^* - \alpha^*/\beta. \quad (49)$$

The cusp line has y-intercept $g_{\max}(\beta)/\beta$. Using the definition $g_{\max}(\beta) = g(\varphi^*, \beta) = \frac{1}{2} \ln(1 - \varphi^{*2}) + \beta \varphi^*$ from (15) together with $\alpha^* = -\frac{1}{2} \ln(1 - \varphi^{*2})$ from (48), the cusp-line intercept becomes

$$\frac{g_{\max}(\beta)}{\beta} = \frac{\frac{1}{2} \ln(1 - \varphi^{*2})}{\beta} + \varphi^* = -\frac{\alpha^*(\beta)}{\beta} + \varphi^* = b_{\text{tangent}}(\beta). \quad (50)$$

The y-intercepts coincide. Two lines with the same slope and the same y-intercept are the same line. The cusp line at sharpness β therefore coincides with the tangent to the blue curve at $(\alpha^*(\beta), \varphi^*(\beta))$ — two independently defined lines, one set by the energy landscape and one by the geometry of $\varphi_{\max}$, turn out to be the same line.

Lower envelope. By the concavity statement at the start of this appendix, every tangent line lies above the blue curve, with equality only at the tangency point. So the cusp line at any sharpness β lies above the blue curve, with equality only at $\alpha = \alpha^*(\beta)$:

$$\varphi_c(\alpha, \beta) \geq \varphi_{\max}(\alpha) \quad \text{for all } \alpha \text{ and all } \beta > 0, \quad (51)$$

with equality at one α for each β . Taking the infimum over β ,

$$\varphi_{\max}(\alpha) = \inf_{\beta > 0} \varphi_c(\alpha, \beta), \quad (52)$$

the lower envelope identity in figure coordinates.

Tangency sweeps the blue curve. The map $\beta \mapsto \varphi^*(\beta)$ is continuous and strictly increasing on $(0, \infty)$, with $\varphi^*(\beta) \rightarrow 0$ as $\beta \rightarrow 0^+$ and $\varphi^*(\beta) \rightarrow 1$ as $\beta \rightarrow \infty$. As β sweeps $(0, \infty)$, the tangency point $(\alpha^*(\beta), \varphi^*(\beta))$ therefore traces the entire blue curve from the origin to its right edge.

Where the Legendre transform sits. The cusp line at sharpness β and the tangent to the blue curve of slope $1/\beta$ are two independently defined objects. The cusp line comes from the energy landscape via the saddle value $g_{\max}(\beta)$; the tangent comes from the geometry of the spherical extreme. The y-intercept of either, expressed as a function of β , is the (concave) Legendre transform of $\varphi_{\max}$:

$$g_{\max}(\beta)/\beta = \sup_{\alpha} [\varphi_{\max}(\alpha) - \alpha/\beta], \quad (53)$$

with the supremum attained at $\alpha = \alpha^*(\beta)$, where the slope-match condition (46) holds. Two lines with the same slope and the same y-intercept are the same line; the Legendre identity therefore expresses the coincidence of the cusp line at β with the tangent line of slope $1/\beta$. The lower-envelope identity (52) is the inverse statement: the spherical extreme is recovered as the lower envelope of the cusp lines parameterised by sharpness β .